

Video Screenshot Tool

Turn a screen-demo recording into timestamped screenshots Claude can pull on demand

What it does. Drop a screen-recording (MP4/MOV) on the app. It transcribes the audio into a **timestamped transcript** — the map of everything said and when. From then on, you just ask Claude for a moment ("*show me where Kristin filters the item codes*") and Claude reads the transcript, finds the timecode, and grabs that **exact** screenshot. No scrubbing, and no wall of redundant images.

The idea in one picture

Step	When	What happens
Transcribe	Up front, automatic	Whisper turns the audio into a transcript with exact timestamps — the searchable map.
Grab a frame	On demand, by Claude	Claude seeks the precise moment you asked about and saves just that screenshot.
Blanket frames	Only if you ask	For silent stretches with no narration, optionally pre-extract frames (see last page).

Frames exist only because they were needed — captured at the exact second, not the nearest sample.

Using it — the everyday flow

- 1 **Drag a recording onto `RedArc Video Tool.app`**. A Terminal window opens and shows progress. The first run downloads the transcription model once (a minute or two); after that it's quick.
- 2 **When it finishes, the output folder opens.** You now have the transcript and a `manifest.json` for that video.
- 3 **Ask Claude for the moment you want.** e.g. "*In the platform-walkthrough video, grab the screen where Zak opens the rate sheet.*" Claude finds the line in the transcript, grabs that frame, and shows it to you.

That's the whole loop. You drop a video and ask in plain English; Claude does the finding and grabbing.

What lands on disk

```
output/<video-name>/
  source.mp4          copy of the original (so grabs work anytime later)
  transcript.txt       readable, [MM:SS]-prefixed
  transcript.json      machine-readable segments + timestamps
  manifest.json        the index Claude reads
  frames/             created on demand
    00-04-22.jpg       ← only the moments that were actually needed
```

Running it by hand (optional)

You never have to, but the same thing works from a Terminal in the tool folder:

```
# 1. transcribe a video
./venv/bin/python ./process_video.py /path/to/demo.mp4

# 2. grab one or more exact frames (seconds, MM:SS, or HH:MM:SS)
./venv/bin/python ./grab.py output/demo 4:22 5:10 6:03
```

When a moment has no narration

On-demand grabbing relies on the transcript describing the moment. If something important happens silently, pre-extract frames so there's something to look at:

```
# a frame on every visual scene change
RAVT_FRAMES=scene ./venv/bin/python ./process_video.py demo.mp4

# scene changes PLUS a guaranteed frame every few seconds
RAVT_FRAMES=both ./venv/bin/python ./process_video.py demo.mp4
```

Settings you can tweak

Setting	Default	Effect
RAVT_FRAMES	none	none (on-demand) · scene · interval · both
RAVT_MODEL	small.en	Transcription model: base.en (faster) or medium.en (more accurate)
RAVT_FLOOR	4	Seconds between forced frames in interval mode
RAVT_SCENE	0.30	Scene sensitivity (0–1); lower catches more changes

First-time notes & troubleshooting

"Unidentified developer" warning. The first time you open the app, macOS may block it. Right-click the app → **Open** → **Open**. You only do this once.

First run is slower. The transcription model downloads once on the very first video, then is cached for every run after.

Moving to another Mac? Install the two free dependencies: `brew install ffmpeg`, then in the tool folder `python3 -m venv venv && ./venv/bin/pip install faster-whisper`. Everything is local — no API keys, no per-video cost.